

CS395T: Sequential CRF for NER Project Report

Zeyuan Hu

Computer Science Department

University of Texas at Austin

Austin, Texas

iamzeyuanhu@utexas.edu

Abstract

In this project, I implement the Hidden Markov Model (HMM) and Conditional Random Field (CRF) to solve the Parts-of-Speech tagging (POS) and the Named Entity Recognition (NER) problems respectively. I extend the NER system with the beam search and the transition potential constraints. A series of experiments show that the implementation of the system and the two extensions can have significant impacts on the system performance.

1 Collaborator

- Zhan Shi zshi17@cs.utexas.edu

2 Introduction

Parts-of-Speech (POS) tagging is a task to give a tag to each word in a given sentence. Tag may include: noun, verb, pronoun, preposition, adverb, and so on. POS is important because knowing the tags of words can give us the information about likely neighbouring words and the syntactic structure of the sentence, which will be useful for Syntactic Parsing, Named Entity Recognition, and other information extraction tasks (Jurafsky and H.Martin, 2017, Chapter 10). However, since the amount of tags and sentences can be large, manually tagging each words in a sentence can be extremely time-consuming. Thus, we want to invoke some statistical model to automatically generate the tags for us, which, in this case, is the Hidden Markov Model (HMM). I model the sequence of words and tags as markov chain and use the maximum likelihood estimation to get the transition and emission probabilities. Lastly, I use the Viterbi algorithm to infer the tags from the words sequence.

Another task I perform is the Named Entity Recognition (NER), which is about finding semantic components for a given sentence. For example, given a sentence: "Barack Obama will travel to Hangzhou today for the G20 meeting.", the NER system will generate the label "PERSON" for "Barack Obama", "LOC" for "Hangzhou" and "ORG" for "G20". The model I use for this task is the Conditional Random Field (CFR), which I directly model the output labels y given the word sequence x . I use the forward-backward algorithm to train the model and use the Viterbi algorithm to perform the inference of the labels.

3 Implementation Details

3.1 HMM

For HMM, I implement the Viterbi algorithm. The code structure follows Figure 10.8 in Jurafsky and H.Martin (2017). There are two major parts: a initialization step and a recursion step. I implement two matrices `viterbi` and `backpointer` with dimensions of N by T , where N represents the number of tags(i.e., states) and T represents the number of words for a given sentence. `viterbi` is used to keep track of the score for a given word and a given tag. `backpointer` is used to store the optimum tags we have visited so far that maximize the score for the current tag given the current word. In the initialization step, I calculate the score of each tag for the first word and initialize the values of `backpointer` to be all zeros. In the recursion step, I recursively calculate the score of each (tag,word) combination and keep track of the optimum paths to each tag of the current word. Once we reach the final word, we can choose the tag that has the highest score for the last word and then construct the best possible tag for each word by recursively backtracking to the first word using the values stored in `backpointer` matrix.

3.2 CRF

For the CRF, I implement the forward-backward algorithm to train the model and the Viterbi algorithm for inference of the labels. The Forward-backward algorithm allows us to compute $P(y_i = s | \mathbf{x})$, which will be used in the Stochastic gradient descent (SGD) to update the weights of the model. I use three matrices `feature_matrix`, `forward`, and `backward` with dimensions of N by T in my implementation. `feature_matrix` is used to hold the calculation result of $\phi_e(y_i, i, \mathbf{x})$. `forward` and `backward` play the same roles as `viterbi` matrix in HMM. The only differences between `forward` and `backward` is how the score gets calculated: `forward` calculates $\alpha_{i,t}$, whereas `backward` calculates $\beta_{i,t}$ (Marsland, 2014). After forward-backward algorithm, I compute the gradient and run the SGD to train the model. The implementation of the Viterbi algorithm in CRF is essentially the same as the one in HMM. The only difference between these two versions is the extra constraints on what labels sequence are considered valid in the CRF. The transition constraints will be described in details in the next section.

3.3 Extensions

My extension to the project is centered around the speed and the accuracy. For the speed, I implement the beam search for both HMM and CRF. For the accuracy, I encode hard constraints in the viterbi algorithm of the CRF so that certain labels sequence will not get picked. For the beam search, I maintain a list of `beam`, which is a priority queue that sorts the scores in the descending order. The implementation looks like the Viterbi algorithm in the way that there are initialization and recursion steps as well. In the initialization step, I set the beam for the first word and keep the labels that have the top `beam_size` scores. In the recursion step, for each label in the previous beam, I expand it using N labels and put the corresponding scores from the calculation into the current beam. At the last, we read through all the top of the beams to recover the best possible labels for the words sequence.

There are three types of labels sequence that are illegal in the NER system: 1. We cannot have "O" label followed by "I" label of any type (i.e., "OI", "I-LOC"). 2. We cannot have two "I" labels with different types next to each other (i.e.,

"I-LOC", "I-PER"). 3. We cannot have "I" label of a type that is different from the "B" label's type that it follows (i.e., "B-LOC", "I-PER"). I handle all these cases by explicitly checking these conditions in the recursion step of the Viterbi algorithm and the beam search.

4 Experiments

Figure 1 shows the result of the F1 score of the NER system on the development set. As the number of epoches in the training phase increases, the F1 scores increase but after around 15 epoches, the F1 score becomes stable and stays above 85. One extension to the system is using the beam search instead of the Viterbi algorithm for the inference. As shown in figure 1, the beam search has the same performance as the Viterbi algorithm in terms of the F1 score. However, the beam search runs faster than the Viterbi algorithm. Profiling of the system shows that the beam search with `beam_size = 2` on average takes 2.18 nanoseconds while the Viterbi algorithm takes 8.21 nanoseconds for the same set of labels and sentences.




Figure 1: F1 score vs. number of epoches in development set

Implementation of the system can have a huge impact on its performance as well. For example, I compare `logsumexp` (Jones et al., 2001–) with `logaddexp` (Jones et al., 2001–) in the forward-backward score calculation. The former supports the vector input but the later can only apply towards two numbers. I adopt the former in my initial implementation and each epoch takes 150 seconds on average. Then, I experiment with the `logaddexp` and find out that each epoch takes 120 seconds on average while keeping the rest im-

plementation untouched. Another example is the learning rate of the SGD. I compare the constant 0.001 learning rate with a rate adjustment scheme of decreasing the rate by 10% per 10 iterations. The experiment shows that constant 0.001 has a very slow convergence to the global optimum: after five iterations, the F1 score is still below 70. On the other hand, the second scheme allows the system's F1 score to pass 80 relatively quickly - around 15 iterations.

Encoding the transition potential constraints into both the Viterbi algorithm and the beam search is the key factor in improving system labelling accuracy. Without the transition potential constraints, the F1 score is capped at 83 no matter the number of epoches used during the training phase. The output of development set also confirms with the F1 scores as many sentences have been mislabeled with illegal labels sequence. After the integration of the transition constraints into the algorithms, the F1 score rises to at least 85 after around 20 epoches.

5 Conclusion and Future Work

In this project, I implement both the HMM and the CRF models for POS tagging and NER respectively. I show the importance of function implementation choice and learning rate towards the NER system performance. I add the beam search as an alternative inference algorithm for both HMM and CRF and the beam search has the same performance as the Viterbi algorithm but with faster execution speed. In the future, the NER system can be further improved by exploring the potential performance gain from the structured SVM instead of the CRF, using the AdaGrad optimizer instead of the SGD, and better feature engineering on the German NER data.

References

- Eric Jones, Travis Oliphant, Pearu Peterson, et al. 2001-. [SciPy: Open source scientific tools for Python](http://www.scipy.org/). [Online; accessed ;today;]. <http://www.scipy.org/>.
- Dan Jurafsky and James H.Martin. 2017. Speech and language processing. Preprint on webpage at <https://web.stanford.edu/~jurafsky/slp3/>.
- Stephen Marsland. 2014. *Machine Learning: An Algorithmic Perspective, Second Edition*. Chapman & Hall/CRC, 2nd edition.