

# CS380L: Advanced Operating Systems Lab #2

Zeyuan Hu <sup>1</sup>, iamzeyuanhu@utexas.edu

EID:zh4378 Spring 2019

## 1 Environment

We have two Linux machines for experiments: `erode-lom` and `thoothukudi-lom`. Both of them run Ubuntu 16.04 LTS. <sup>2</sup>

## 2 User-level File System

### 2.1 Implementation

We leverage `libfuse` and `libssh` to implement a network file system. Filesystem in Userspace (FUSE) is a user-space file system framework. FUSE consists of a Linux kernel module and a user-level daemon. When a user application performs operations on a mounted FUSE file system, the operation will be routed to FUSE's kernel driver by VFS. The operations as requests will be maintained by a queue and user-level daemon will pick a request from the kernel queue and process the request. Daemon will write response back to kernel once it is done with processing. More information about FUSE can be seen in [1].

#### 2.1.1 System Architecture

We implement a network file system. Like NFS, it supports a server with multiple clients. On the client side, the FUSE client program will mount a user-space file system with remote user, remote host, remote path, local cache path, and local mount point provided by the user. Remote user and remote host specify the server that our network file system want to contact. Remote path specifies the location on the server we want to store and fetch files for the clients. Our system requires user to specify a path for local cache as we serve clients' requests from local cache as much as we can. In other words, some file operations are directly served by the local cache without further contact with server. Communications between clients and the server are done via SSH protocol. Thus, there is no server-side implementation in our system.

---

<sup>1</sup>30 hours spent on this lab.

<sup>2</sup>`erode-lom` with ip address: 192.168.1.120 and `thoothukudi-lom` with ip address: 192.168.1.121. Due to the network setup, those two machines are not properly registered in the department DNS server. We use ip address instead host names to logon these two machines.

### 2.1.2 FUSE Calls Implemented

We implement the following operations: `getattr`, `readdir`, `create`, `open`, `read`, `write`, `release`, `fsync`, `mkdir`, `unlink`, and `rmdir` in our system. Our operation implementation is based on `libfuse` version is 2.9.4 <sup>3</sup> API. In `libfuse`, there are two sets of APIs: a "high-level", synchronous API, and a "low-level" asynchronous API. The key difference between "high-level" and "low-level" API is that "high-level" allows us to work with file names and path instead of inodes in synchronous fashion. Thus, for the simplicity, our implementation adopts the "high-level" API.

`getattr` operation is used to get attributes of a file or a directory. Function signature of `libfuse`'s `getattr` operation is `int getattr(const char * path, struct stat * stbuf);`. Our implementation should fill `stbuf` to contain attributes of file or directory indicated by `path`. Since `struct stat` is the same structure used in `stat` call [2], we can leverage `stat` command [3]. Specifically, in our implementation, we execute `stat` command remotely via SSH to collect attributes of the desired target (specified by `path`). The output of the `stat` command is parsed on the client side of the file system. We fill `st_size`, `st_blocks`, `st_mode`, and `st_nlink` fields of `struct stat` from the parse result. `st_uid` and `st_gid` are filled with UID and GID of the file system client process. `st_mode` is filled with 644 (i.e., `rw-r--r--`) if the target is file and 755 (i.e., `rxr-xr-x`) if the target is directory.

## 3 System Tools Exercise

### 3.1 strace

`script` command allows user to record terminal printout into a file [4]. Per the lab instruction, we use `strace` to trace the syscalls and signals of a target process [5]. In our case, we trace the process involving `cat`. One thing I notice is that `script` contains some unicode as shown in Figure ?? . Thus, we use the following code to clean up the output:

```
1 cat $FILE | perl -pe 's/\e([^\[\]]|\.[.*?[a-zA-Z]|\.[.*?\a])//g' | col -b > $FILE-processed
```

The result is shown in Figure 1 <sup>4</sup>.

### 3.2 lsof

`lsof` lists all open files [6]. `lsof | grep /dev` shows all the open devices used by user-space programs. On our machine, we have the following opened devices:

---

<sup>3</sup>check via `fusermount -V`

<sup>4</sup>raw output and cleanup output comes with the report as `session.record` and `session.record-processed` respectively

```

Script started on Fri 01 Mar 2019 04:37:34 PM CST
ESC[1;34m zeyuanhu @ ESC[0;36mHotDog(thoothukudi-lom)ESC[0mESC[1;34m ~ESC[0m
ESC[0;36m Fri Mar 01 16:37:34 $ ESC[0;39mstrace cat - > new_file
execve("/bin/cat", ["cat", "-"], [/ 37 vars */]) = 0
brk(NULL)                                = 0x2280000
access("/etc/ld.so.nohwcap", F_OK)        = -1 ENOENT (No such file or directory)
access("/etc/ld.so.preload", R_OK)        = -1 ENOENT (No such file or directory)

```

Raw output of script, which contains Unicode character.

```

Script started on Fri 01 Mar 2019 04:37:34 PM CST
zeyuanhu @ HotDog(thoothukudi-lom) ~
Fri Mar 01 16:37:34 $ strace cat - > new_file
execve("/bin/cat", ["cat", "-"], [/ 37 vars */]) = 0
brk(NULL)                                = 0x2280000
access("/etc/ld.so.nohwcap", F_OK)        = -1 ENOENT (No such file or directory)
access("/etc/ld.so.preload", R_OK)        = -1 ENOENT (No such file or directory)

```

Script output after cleanup

Figure 1: script output before and after cleanup

- /dev/null: null device
- /dev/pts/\* and /dev/tty: terminal devices
- /dev/urandom: kernel random number source device
- /dev/ptmx: a character file to create a pseudoterminal master

## 4 Network Tools

`ifconfig` command lists all the network interfaces the machine is using to communicate externally. On our machine, interface for Ethernet is `eno1`. We can find IP address, gateway address, and subnet mask from the output.

`tcpdump` command can dump traffic on a network interface. We use the `tcpdump` output provided by the lab instruction to answer the questions below.

### a. Are DHCP messages sent over UDP or TCP?

We use `tcpdump -nn -r tcpdump.out.1 | grep -i dhcp` to filter out the DHCP messages from the dump. `-nn` ensures that we can see the actual port number instead of the port name. One line of the output is:

```

10:19:24.525962 IP 0.0.0.0.68 > 255.255.255.255.67: BOOTP/DHCP, Request from
a8:20:66:3b:66:51, length 300

```

The first field shows the time that the packet was traveling. The second field shows the source host address and port, followed by the destination host address and port. The third field

shows the protocol the packet was using. From DHCP [7], we know DHCP messages sent over UDP. As shown by the printout, the messages are sent between port 68 (client) and port 67 (server).

**b. What is the link-layer (e.g., Ethernet) address of your host? (Feel free to obscure the last couple bytes for privacy's sake)**

We use the same `tcpdump` command as above with extra `-e` option to show link-layer header. The following printout contains DHCP messages for acquiring IP address:

```
10:19:24.525962 a8:20:66:3b:66:51 > ff:ff:ff:ff:ff:ff, ethertype IPv4 (0x0800
), length 342: 0.0.0.0.68 > 255.255.255.255.67: BOOTP/DHCP, Request from
a8:20:66:3b:66:51, length 300
10:19:24.566258 00:21:9b:fb:61:0c > a8:20:66:3b:66:51, ethertype IPv4 (0x0800
), length 342: 128.83.158.2.67 > 128.83.158.160.68: BOOTP/DHCP, Reply,
length 30
```

From the printout we can see that the link-layer address (MAC address) of the host is a8:20:66:3b:66:51.

**c. What is the IP address of your DHCP server?**

From the printout above, we can see the IP address of DHCP server is 128.83.158.2 and the new IP address acquired from DHCP server is 128.83.158.160.

**d. What is the purpose of the DHCP release message?**

DHCP release message is used to release IP address.

**e. Does the DHCP server issue an acknowledgment of receipt of the client's DHCP request?**

DHCP server does not issue an acknowledgment of receipt of the client's release message.

**f. What would happen if the client's DHCP release message is lost?**

If DHCP release message is lost, the DHCP server has to wait for the lease to timeout before assigning it to other clients.

## References

- [1] B. K. R. Vangoor, V. Tarasov, and E. Zadok, "To {FUSE} or not to {FUSE}: Performance of user-space file systems," in *15th {USENIX} Conference on File and Storage Technologies ({FAST} 17)*, pp. 59–72, 2017.

- [2] “stat(2) - linux man page.” <https://linux.die.net/man/2/stat>.
- [3] “stat(1) - linux man page.” <https://linux.die.net/man/1/stat>.
- [4] “script(1) - linux man page.” <http://man7.org/linux/man-pages/man1/script.1.html>.
- [5] “strace(1) - linux man page.” <https://linux.die.net/man/1/strace>.
- [6] “lsof(8) - linux man page.” <http://man7.org/linux/man-pages/man8/lsof.8.html>.
- [7] “Dynamic host configuration protocol.” [https://en.wikipedia.org/wiki/Dynamic\\_Host\\_Configuration\\_Protocol](https://en.wikipedia.org/wiki/Dynamic_Host_Configuration_Protocol), 2018.